



## Module 06 – Exercise Class

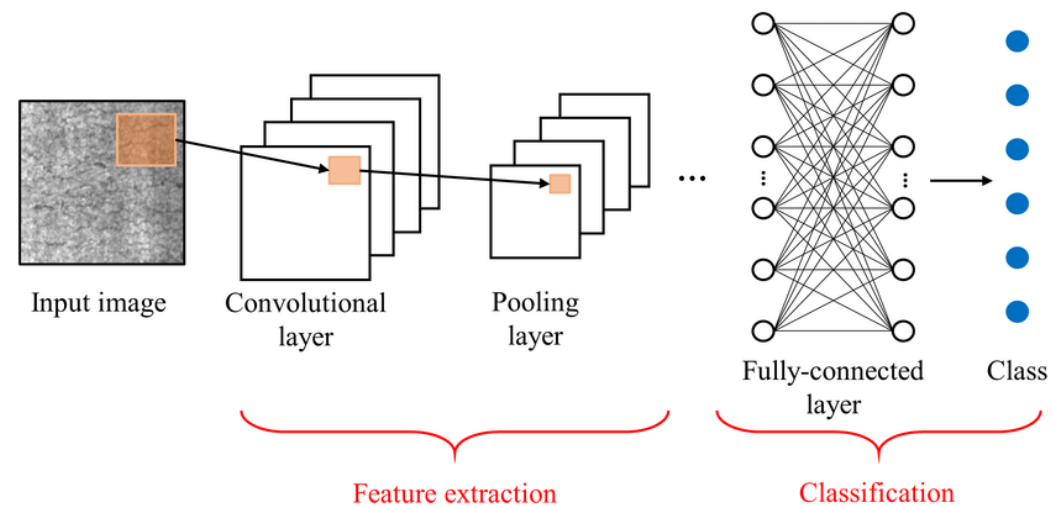

# Convolutional Neural Network

Nguyen Quoc Thai

# Objectives

## CNN

- ❖ Convolution Layer, Pooling Layer
- ❖ Multiple Input – Output Channels
- ❖ LeNet Model




## Classification

- ❖ Image Classification
- ❖ TextCNN Model
- ❖ Text Classification

# Outline

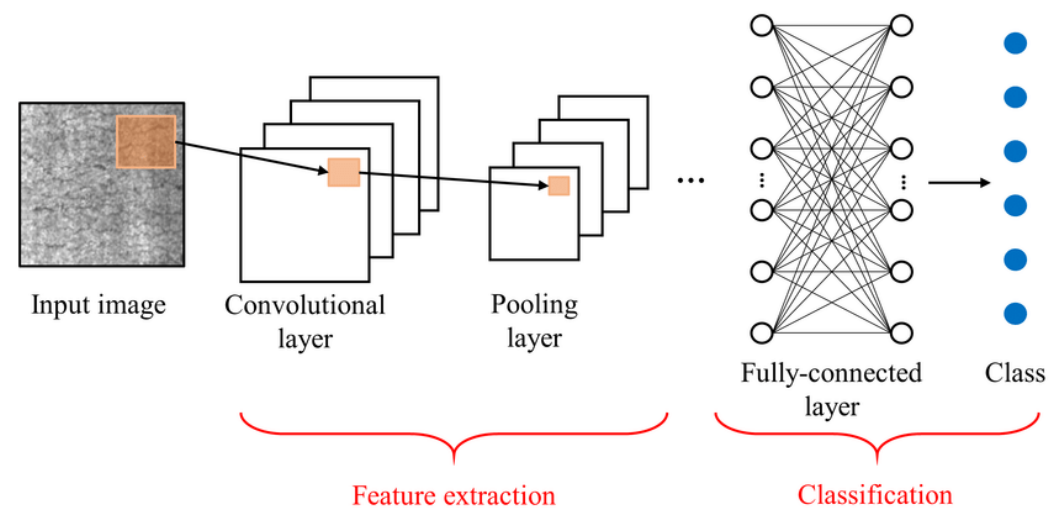
## SECTION 1

### CNN



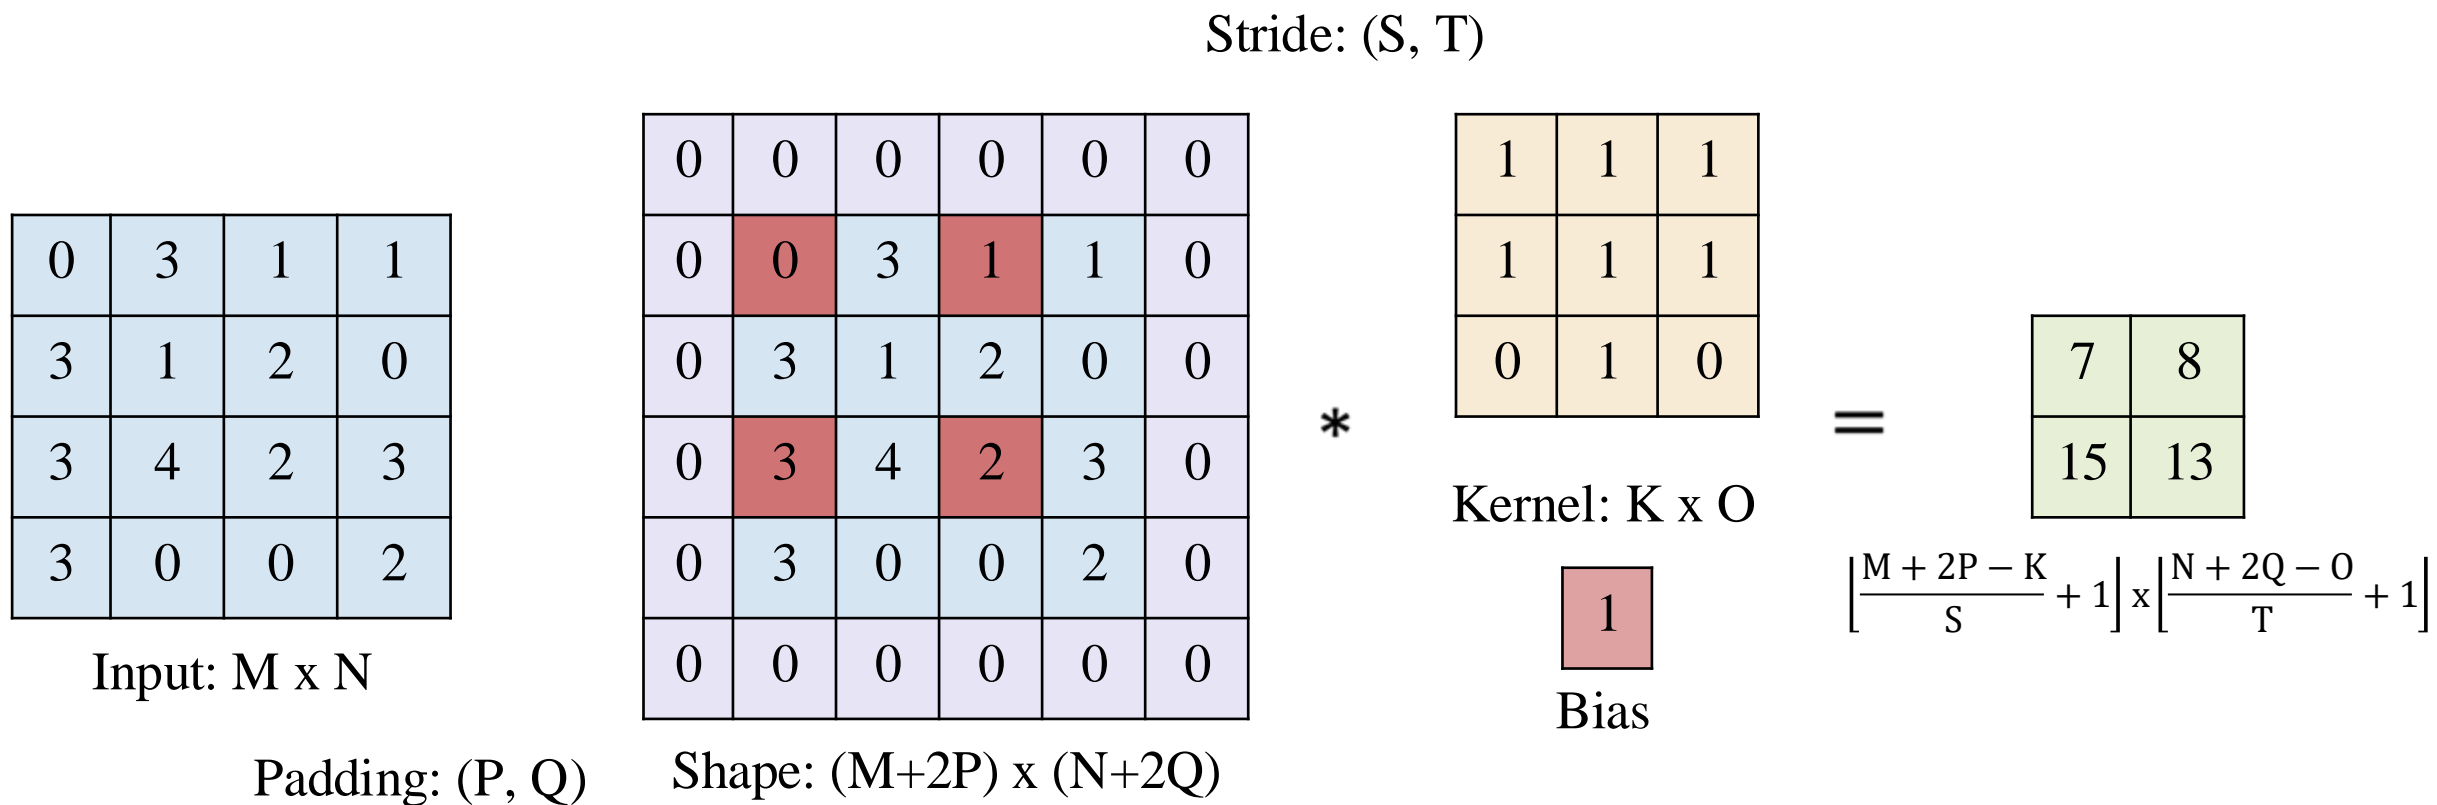
## SECTION 2

### Applications





## Convolutional Layer





## Max Pooling Layer

3	2	1	0	0	3
0	3	3	1	1	0
3	1	4	1	1	0
2	4	1	1	0	4
1	0	3	0	3	0
3	4	4	3	3	4

Input: 6 x 6

Kernel Size: 2  
Stride: 2

3	3	3
4	4	4
4	4	4

Output: 3 x 3

## MaxPool1d

Kernel Size: 3

Stride: 3

3	3
3	1
4	1
4	4
3	3
4	4

Output: 6 x 2



## Average Pooling Layer

3	2	1	0	0	3
0	3	3	1	1	0
3	1	4	1	1	0
2	4	1	1	0	4
1	0	3	0	3	0
3	4	4	3	3	4

Input: 6 x 6

Kernel Size: (3, 2)  
Stride: 2

2.0	1.7	0.8
1.8	1.6	1.3

Output: 2 x 3

**AvgPool1d**  
Kernel Size: 3  
Stride: 3

2.0	1.0
2.0	0.7
2.7	0.7
2.3	1.7
1.3	1.0
3.7	3.3

Output: 6 x 2



## Multiple Inout Channels

0	0	2	1
4	1	3	1
4	3	1	4
2	4	2	0

Input Channel #1 (Red)

0	4	3	4
4	1	2	0
1	2	2	4
2	3	3	4

Input Channel #2 (Green)

3	2	0	4
0	4	1	4
0	4	3	3
1	0	0	0

Input Channel #3 (Blue)

1	1	1
1	1	1
1	0	0

Kernel Channel #1

14

+

1	1	1
1	1	1
1	1	0

Kernel Channel #2

17

+

1	0	1
1	0	1
0	1	1

Kernel Channel #3

11

+

Bias

1
---

43	50
40	13



## Multiple Input Channels

```
input = torch.randint(5, (3, 4, 4), dtype=torch.float32)
input
```

```
tensor([[[0., 0., 2., 1.],
         [4., 1., 3., 1.],
         [4., 3., 1., 4.],
         [2., 4., 2., 0.]],
```

```
        [[0., 4., 3., 4.],
         [4., 1., 2., 0.],
         [1., 2., 2., 4.],
         [2., 3., 3., 4.]],
```

```
        [[3., 2., 0., 4.],
         [0., 4., 1., 4.],
         [0., 4., 3., 3.],
         [1., 0., 0., 0.]])])
```

```
# init weight
```

```
conv_layer.weight.data = init_kernel_weight
conv_layer.weight
```

```
Parameter containing:
```

```
tensor([[[[1., 1., 1.],
          [1., 1., 1.],
          [1., 0., 0.]],
```

```
        [[1., 1., 1.],
         [1., 1., 1.],
         [1., 1., 0.]],
```

```
        [[1., 0., 1.],
         [1., 0., 1.],
         [0., 1., 1.]]]), requires_grad=True)
```

```
# define convolutional layer
```

```
conv_layer = nn.Conv2d(
    in_channels=3,
    out_channels=1,
    kernel_size=3,
)
```

```
# init bias
```

```
conv_layer.bias = nn.Parameter(
    torch.tensor([1], dtype=torch.float32)
)
conv_layer.bias
```

```
Parameter containing:
```

```
tensor([1.], requires_grad=True)
```

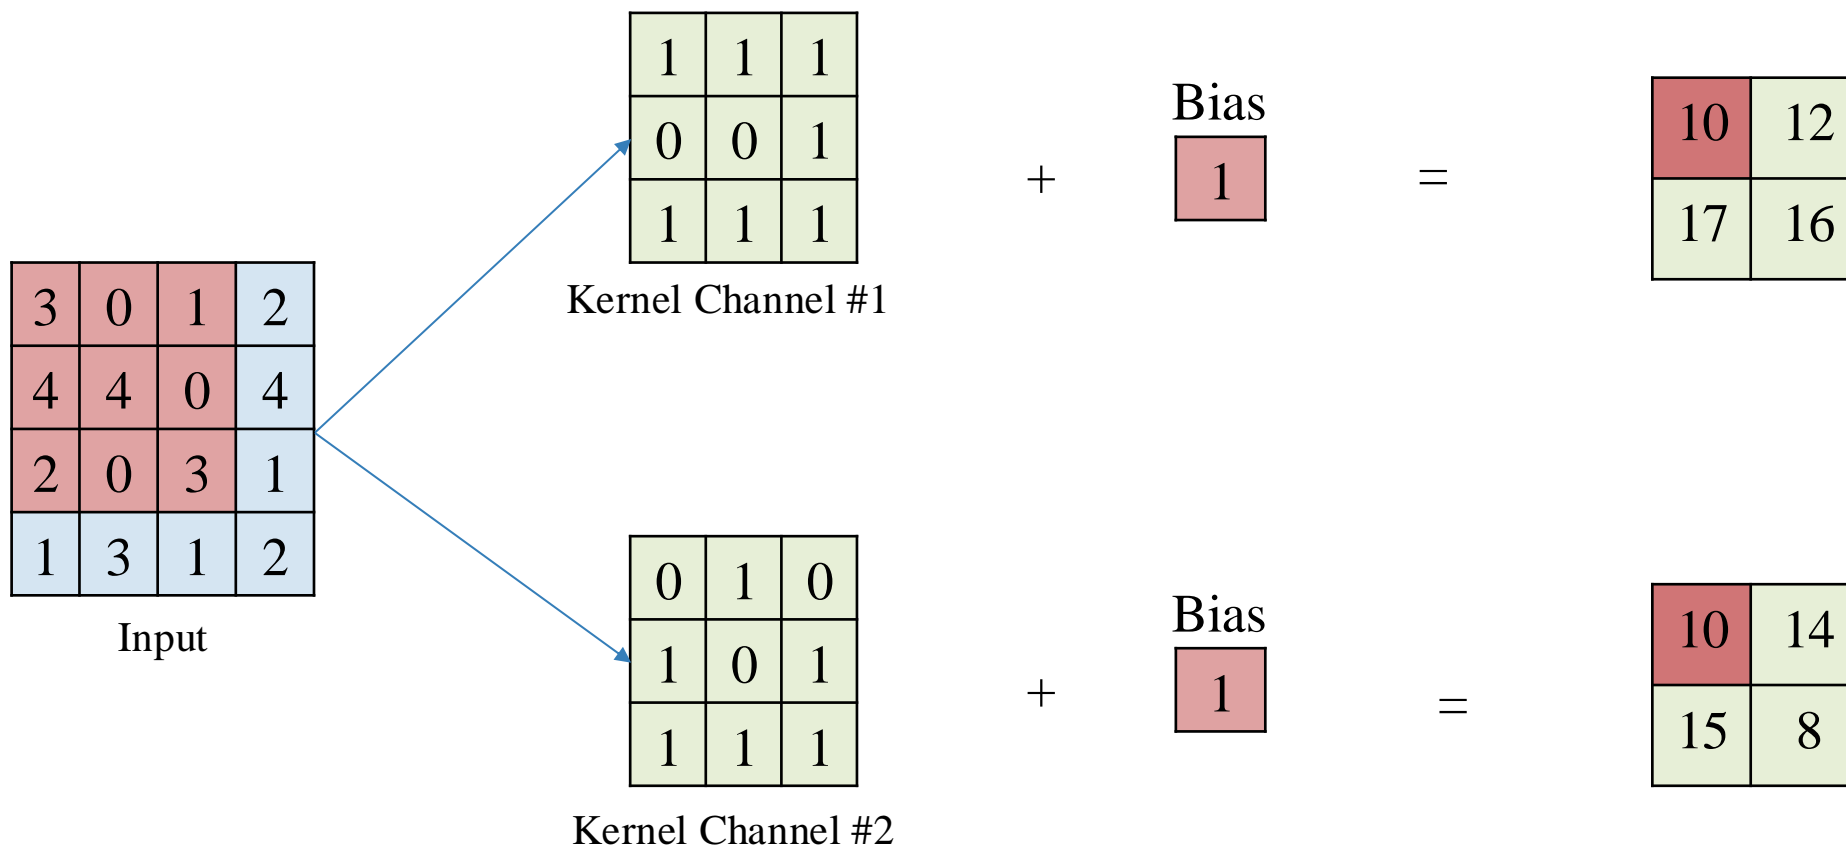
```
output = conv_layer(input)
output
```

```
tensor([[[43., 50.],
         [40., 50.]]], grad_fn=<SqueezeBackward1>)
```





## Multiple Output Channels





## Multiple Output Channels

```
input = torch.randint(5, (1, 4, 4), dtype=torch.float32)
input
```

```
tensor([[[[3., 0., 1., 2.],
          [4., 4., 0., 4.],
          [2., 0., 3., 1.],
          [1., 3., 1., 2.]]]])
```

```
# init weight
```

```
conv_layer.weight.data = init_kernel_weight
conv_layer.weight
```

```
Parameter containing:
```

```
tensor([[[[1., 1., 1.],
          [0., 0., 1.],
          [1., 1., 1.]]]],
```

```
        [[0., 1., 0.],
         [1., 0., 1.],
         [1., 1., 1.]]]], requires_grad=True)
```

```
input = torch.randint(5, (1, 4, 4), dtype=torch.float32)
input
```

```
tensor([[[[3., 0., 1., 2.],
          [4., 4., 0., 4.],
          [2., 0., 3., 1.],
          [1., 3., 1., 2.]]]])
```

```
# init bias
```

```
conv_layer.bias = nn.Parameter(
    torch.tensor([1, 1], dtype=torch.float32)
)
conv_layer.bias
```

```
Parameter containing:
```

```
tensor([1., 1.], requires_grad=True)
```

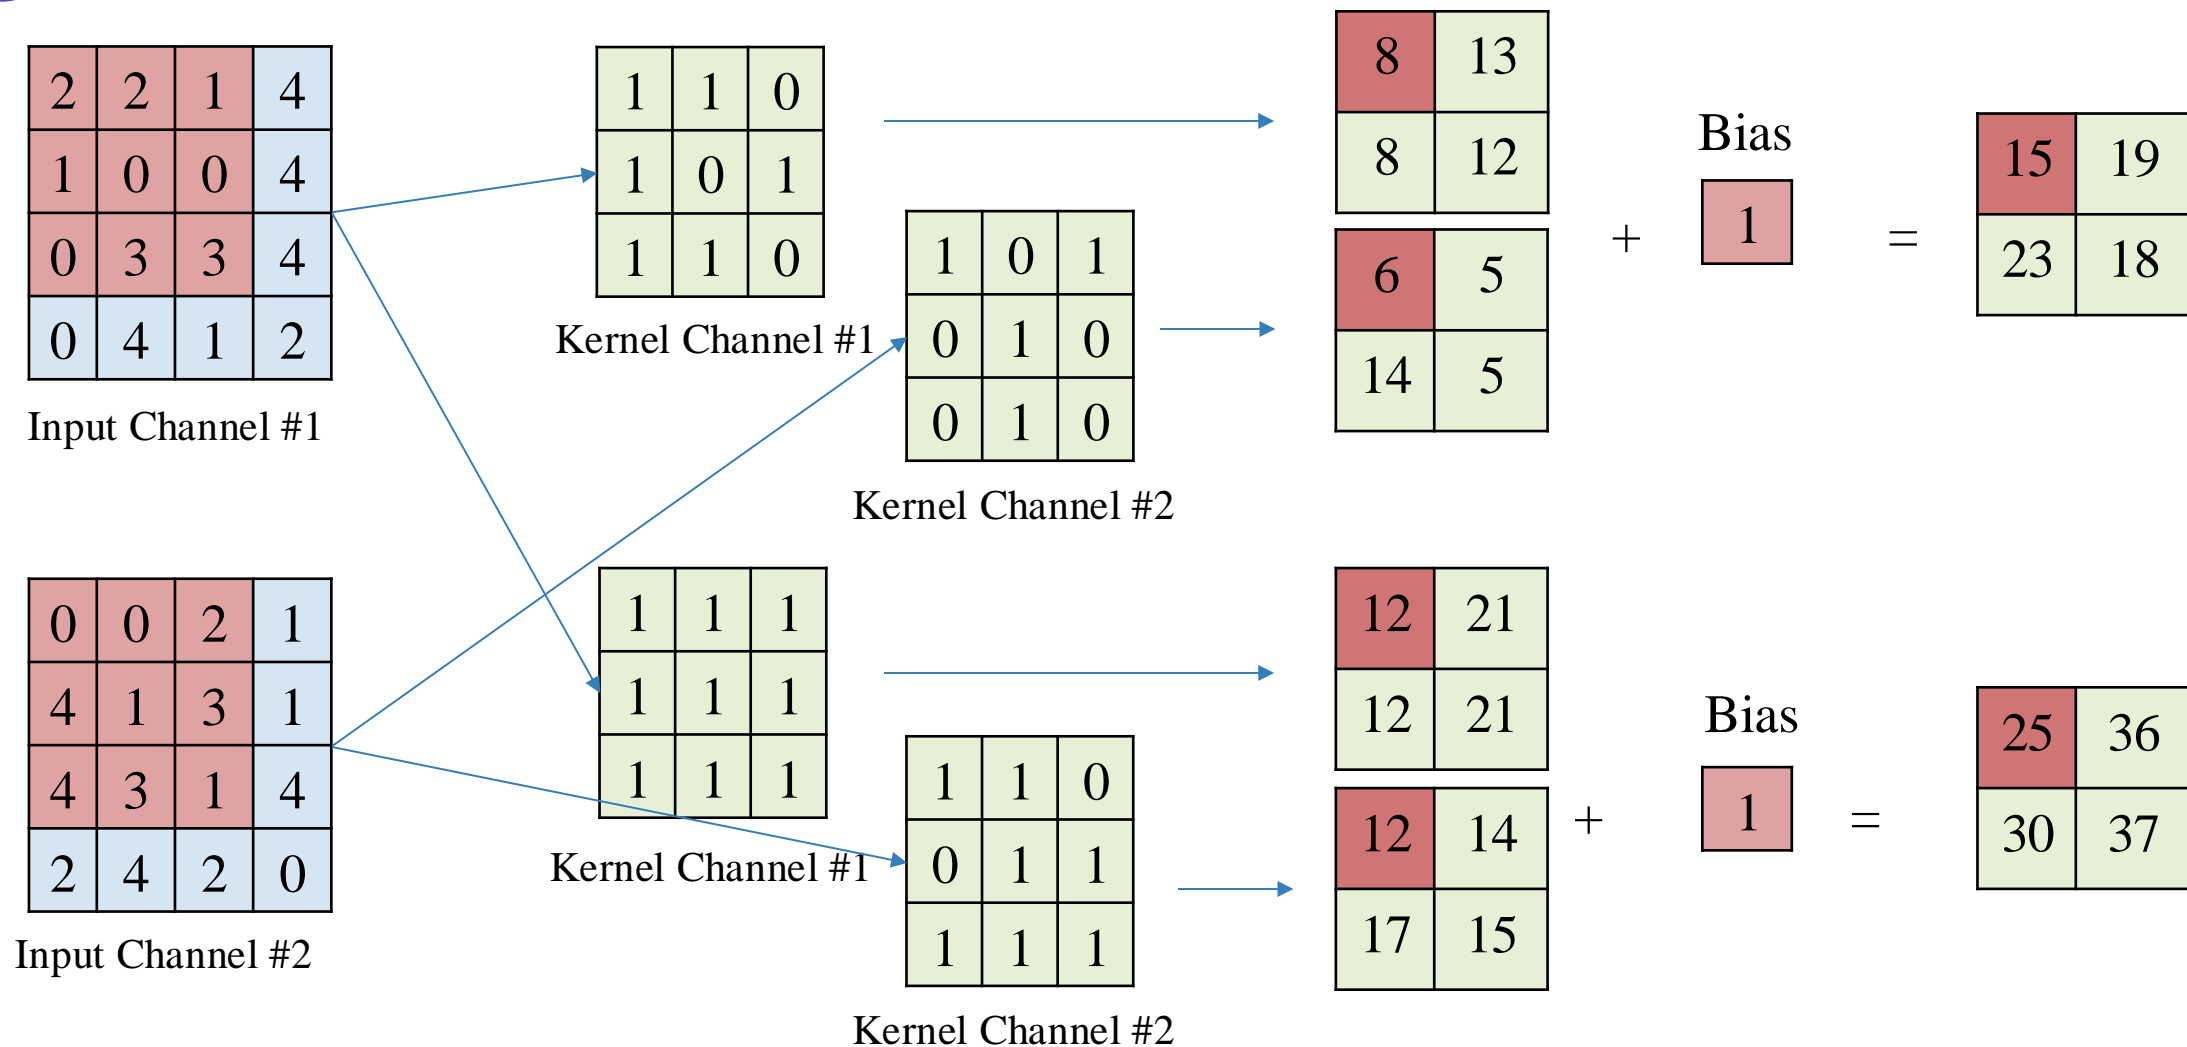
```
output = conv_layer(input)
output
```

```
tensor([[[[10., 12.],
          [17., 16.]],
```

```
        [[10., 14.],
         [15., 8.]]], grad_fn=<SqueezeBackward1>)
```



## Multiple Input – Output Channels





## Multiple Input – Output Channels

```
input = torch.randint(5, (2, 4, 4), dtype=torch.float32)
input
```

```
tensor([[[2., 2., 1., 4.],
         [1., 0., 0., 4.],
         [0., 3., 3., 4.],
         [0., 4., 1., 2.]],
```

```
        [[0., 0., 2., 1.],
         [4., 1., 3., 1.],
         [4., 3., 1., 4.],
         [2., 4., 2., 0.]])
```

```
# define convolutional layer
conv_layer = nn.Conv2d(
    in_channels=2,
    out_channels=2,
    kernel_size=3,
)
```

```
# init weight
conv_layer.weight.data = init_kernel_weight
conv_layer.weight
```

```
Parameter containing:
tensor([[[[1., 1., 0.],
         [1., 0., 1.],
         [1., 1., 0.]],
```

```
        [[1., 0., 1.],
         [0., 1., 0.],
         [0., 1., 0.]]],
```

```
        [[1., 1., 1.],
         [1., 1., 1.],
         [1., 1., 1.]],
```

```
        [[1., 1., 0.],
         [0., 1., 1.],
         [1., 1., 1.]]]), requires_grad=True
```

```
# init bias
conv_layer.bias = nn.Parameter(
    torch.tensor([1, 1], dtype=torch.float32)
)
conv_layer.bias
```

```
Parameter containing:
tensor([1., 1.], requires_grad=True)
```

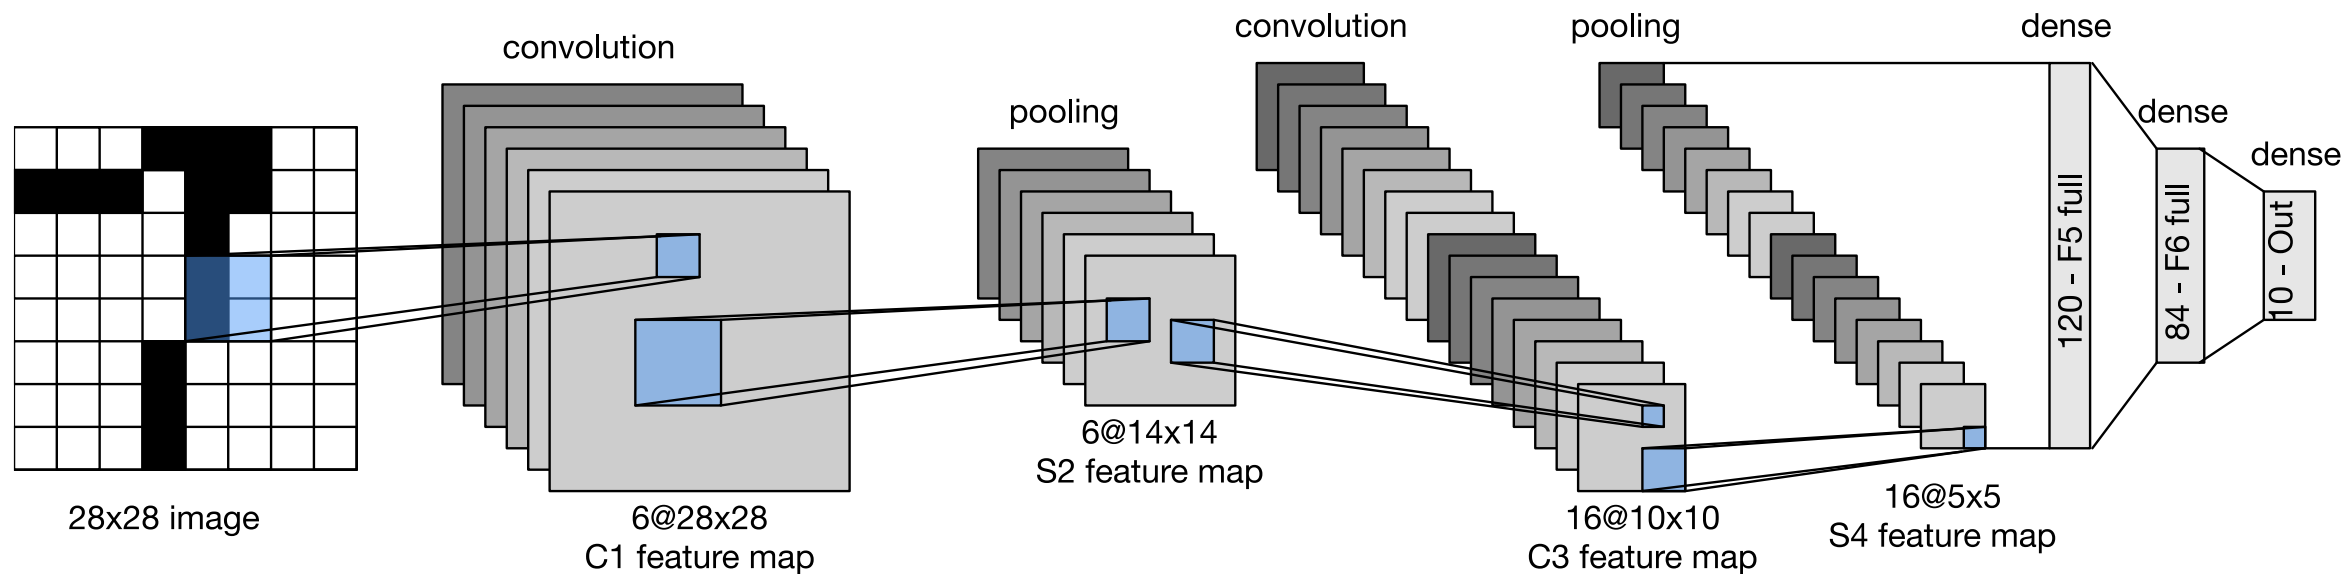
```
output = conv_layer(input)
output
```

```
tensor([[[15., 19.],
         [23., 18.]],
```

```
        [[25., 36.],
         [30., 37.]]], grad_fn=<SqueezeBackward1>)
```



## LeNet Model





## LeNet Model

```
1 class LeNetClassifier(nn.Module):
2     def __init__(self, num_classes):
3         super().__init__()
4         self.conv1 = nn.Conv2d(
5             in_channels=1, out_channels=6, kernel_size=5, padding='same')
6         self.avgpool1 = nn.AvgPool2d(kernel_size=2)
7         self.conv2 = nn.Conv2d(in_channels=6, out_channels=16, kernel_size=5)
8         self.avgpool2 = nn.AvgPool2d(kernel_size=2)
9         self.flatten = nn.Flatten()
10        self.fc_1 = nn.Linear(16 * 5 * 5, 120)
11        self.fc_2 = nn.Linear(120, 84)
12        self.fc_3 = nn.Linear(84, num_classes)
13
14    def forward(self, inputs):
15        outputs = self.conv1(inputs)
16        outputs = self.avgpool1(outputs)
17        outputs = F.relu(outputs)
18        outputs = self.conv2(outputs)
19        outputs = self.avgpool2(outputs)
20        outputs = F.relu(outputs)
21        outputs = self.flatten(outputs)
22        outputs = self.fc_1(outputs)
23        outputs = self.fc_2(outputs)
24        outputs = self.fc_3(outputs)
25        return outputs
```



# Outline

## SECTION 1

**CNN**

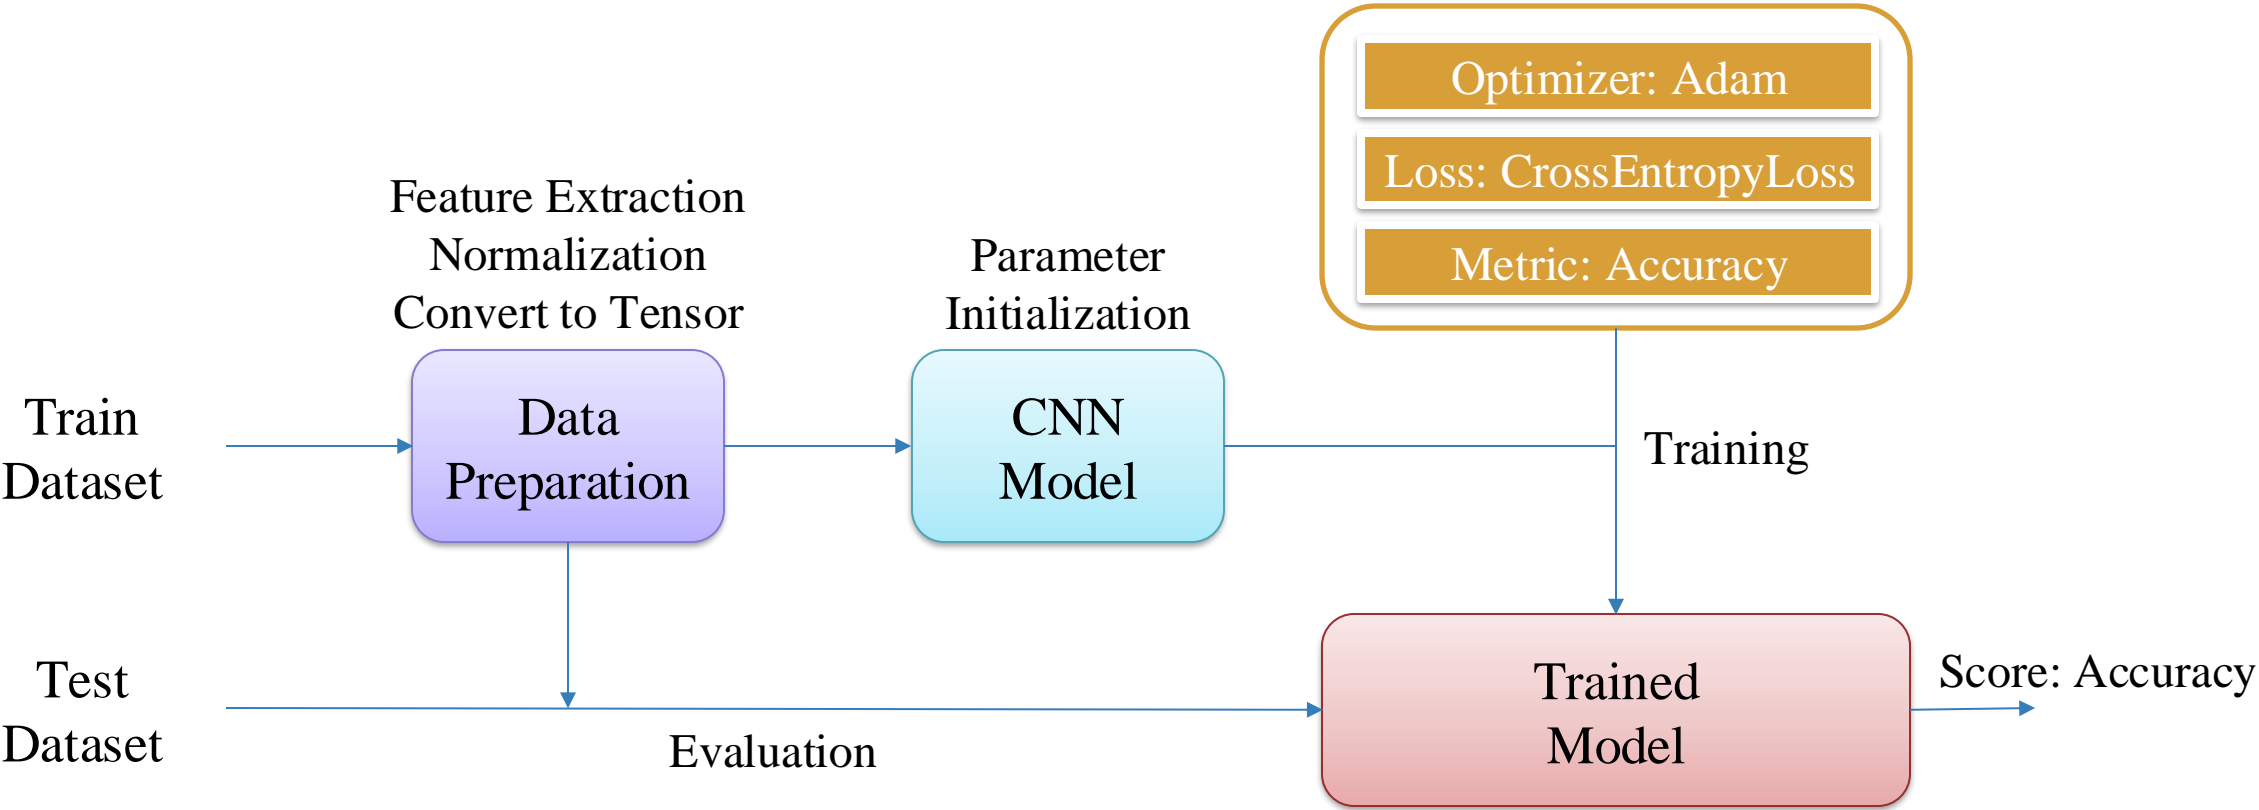
## SECTION 2

**Application**

# Classification



## Pipeline






# Image Classification



## MNIST Dataset

- Images: 70.000
- Class: 10
- Image Size: 28 x 28



# Image Classification



## MNIST Dataset

### ➤ Load dataset

```
ROOT = './data'

train_data = datasets.MNIST(
    root=ROOT,
    train=True,
    download=True
)

test_data = datasets.MNIST(
    root=ROOT,
    train=False,
    download=True
)
```

```
train_data.classes
```

```
['0 - zero',
 '1 - one',
 '2 - two',
 '3 - three',
 '4 - four',
 '5 - five',
 '6 - six',
 '7 - seven',
 '8 - eight',
 '9 - nine']
```

# Image Classification



## MNIST Dataset

### ➤ Preprocessing

```
VALID_RATIO = 0.9
```

```
n_train_examples = int(len(train_data) * VALID_RATIO)
n_valid_examples = len(train_data) - n_train_examples
```

```
train_data, valid_data = data.random_split(
    train_data,
    [n_train_examples, n_valid_examples]
)
```

```
# compute mean and std
mean = train_data.dataset.data.float().mean() / 255
std = train_data.dataset.data.float().std() / 255
mean, std
```

```
(tensor(0.1307), tensor(0.3081))
```

```
train_transforms = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize(mean=[mean], std=[std])
])
```

```
test_transforms = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize(mean=[mean], std=[std])
])
```

```
train_data.dataset.transform = train_transforms
valid_data.dataset.transform = test_transforms
```

# Image Classification



## MNIST Dataset

### ➤ Model



```
1 class LeNetClassifier(nn.Module):
2     def __init__(self, num_classes):
3         super().__init__()
4         self.conv1 = nn.Conv2d(
5             in_channels=1, out_channels=6, kernel_size=5, padding='same')
6         self.avgpool1 = nn.AvgPool2d(kernel_size=2)
7         self.conv2 = nn.Conv2d(in_channels=6, out_channels=16, kernel_size=5)
8         self.avgpool2 = nn.AvgPool2d(kernel_size=2)
9         self.flatten = nn.Flatten()
10        self.fc_1 = nn.Linear(16 * 5 * 5, 120)
11        self.fc_2 = nn.Linear(120, 84)
12        self.fc_3 = nn.Linear(84, num_classes)
13
14    def forward(self, inputs):
15        outputs = self.conv1(inputs)
16        outputs = self.avgpool1(outputs)
17        outputs = F.relu(outputs)
18        outputs = self.conv2(outputs)
19        outputs = self.avgpool2(outputs)
20        outputs = F.relu(outputs)
21        outputs = self.flatten(outputs)
22        outputs = self.fc_1(outputs)
23        outputs = self.fc_2(outputs)
24        outputs = self.fc_3(outputs)
25        return outputs
```

# Image Classification



## MNIST Dataset

➤ Training (Test Set: 98%)

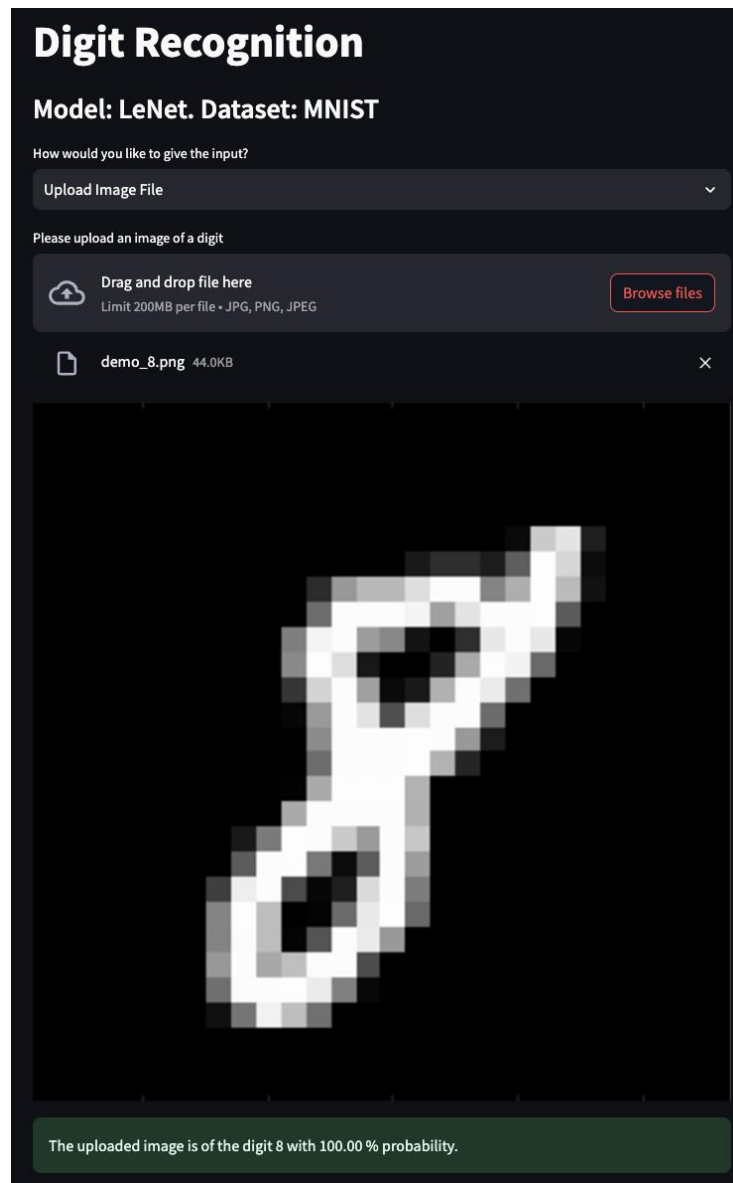


# Image Classification



## MNIST Dataset

➤ Deployment (Streamlit)



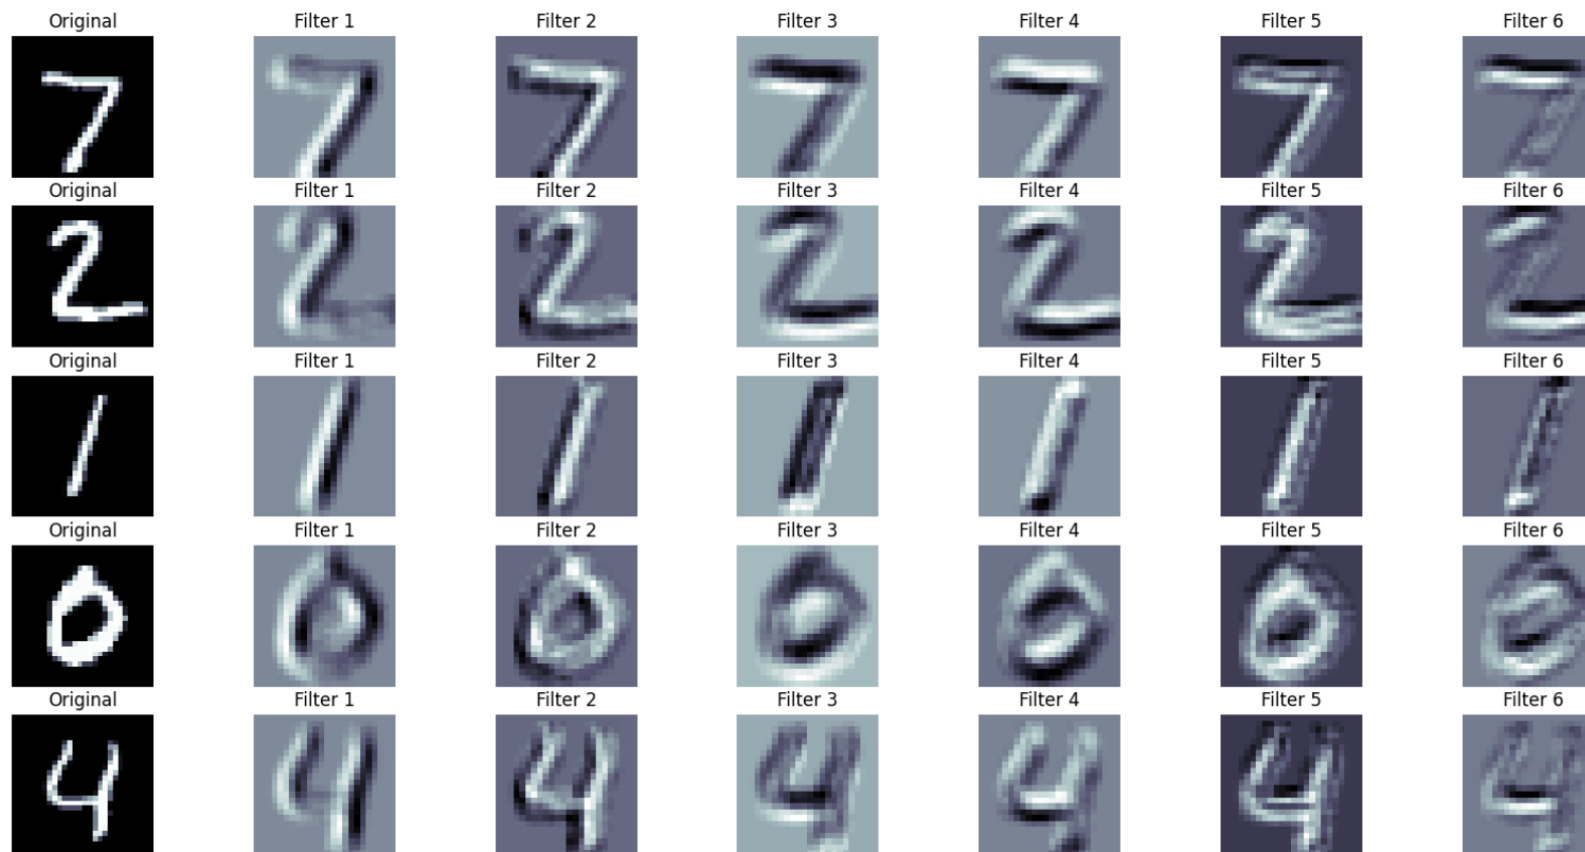
QUIZ TIME

# Image Classification



## MNIST Dataset

### ➤ Kernel Visualization






# Image Classification



## Cassava Leaf Disease Dataset

➤ Class: 5



# Image Classification



## Cassava Leaf Disease Dataset

### ➤ Preprocessing

- ▼ cassavaleafdata
  - ▼ test
    - cbb
    - cbsd
    - cgm
    - cmd
    - healthy
  - ▼ train
    - cbb
    - cbsd
    - cgm
    - cmd
    - healthy
  - ▼ validation
    - cbb
    - cbsd
    - cgm
    - cmd
    - healthy

```
# load image from path
def loader(path):
    return Image.open(path)
```

```
img_size = 150
```

```
train_transforms = transforms.Compose([
    transforms.Resize((150, 150)),
    transforms.ToTensor(),
])
```

```
train_data = datasets.ImageFolder(
    root=data_paths['train'],
    loader=loader,
    transform=train_transforms
)
valid_data = datasets.ImageFolder(
    root=data_paths['valid'],
    transform=train_transforms
)
test_data = datasets.ImageFolder(
    root=data_paths['test'],
    transform=train_transforms
)
```

# Image Classification



## Cassava Leaf Disease Dataset

### ➤ Model

```
summary(lenet_model, (3, 150, 150))
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 6, 150, 150]	456
AvgPool2d-2	[-1, 6, 75, 75]	0
Conv2d-3	[-1, 16, 71, 71]	2,416
AvgPool2d-4	[-1, 16, 35, 35]	0
Flatten-5	[-1, 19600]	0
Linear-6	[-1, 120]	2,352,120
Linear-7	[-1, 84]	10,164
Linear-8	[-1, 5]	425

Total params: 2,365,581

Trainable params: 2,365,581

Non-trainable params: 0

Input size (MB): 0.26

Forward/backward pass size (MB): 2.20

Params size (MB): 9.02

Estimated Total Size (MB): 11.48

```
class LeNetClassifier(nn.Module):
    def __init__(self, num_classes):
        super().__init__()
        self.conv1 = nn.Conv2d(
            in_channels=3, out_channels=6, kernel_size=5, padding='same'
        )
        self.avgpool1 = nn.AvgPool2d(kernel_size=2)
        self.conv2 = nn.Conv2d(in_channels=6, out_channels=16, kernel_size=5)
        self.avgpool2 = nn.AvgPool2d(kernel_size=2)
        self.flatten = nn.Flatten()
        self.fc_1 = nn.Linear(16 * 35 * 35, 120)
        self.fc_2 = nn.Linear(120, 84)
        self.fc_3 = nn.Linear(84, num_classes)



    def forward(self, inputs):
        outputs = self.conv1(inputs)
        outputs = self.avgpool1(outputs)
        outputs = F.relu(outputs)
        outputs = self.conv2(outputs)
        outputs = self.avgpool2(outputs)
        outputs = F.relu(outputs)
        outputs = self.flatten(outputs)
        outputs = self.fc_1(outputs)
        outputs = self.fc_2(outputs)
        outputs = self.fc_3(outputs)
        return outputs
```

# Image Classification



## Cassava Leaf Disease Dataset

➤ Training (Test Set: 53%)



# Image Classification



## Cassava Leaf Disease Dataset

➤ Deployment (Streamlit)

### Cassava Leaf Disease Classification

Model: LeNet. Dataset: Cassava Leaf Disease

How would you like to give the input?

Upload Image File

Please upload an image



Drag and drop file here

Limit 200MB per file • JPG, PNG, JPEG

Browse files



demo\_cbsd.jpg 95.1KB



The uploaded image is of the cbsd with 45.78 % probability.

# Text Classification



## NTC-SCV Dataset

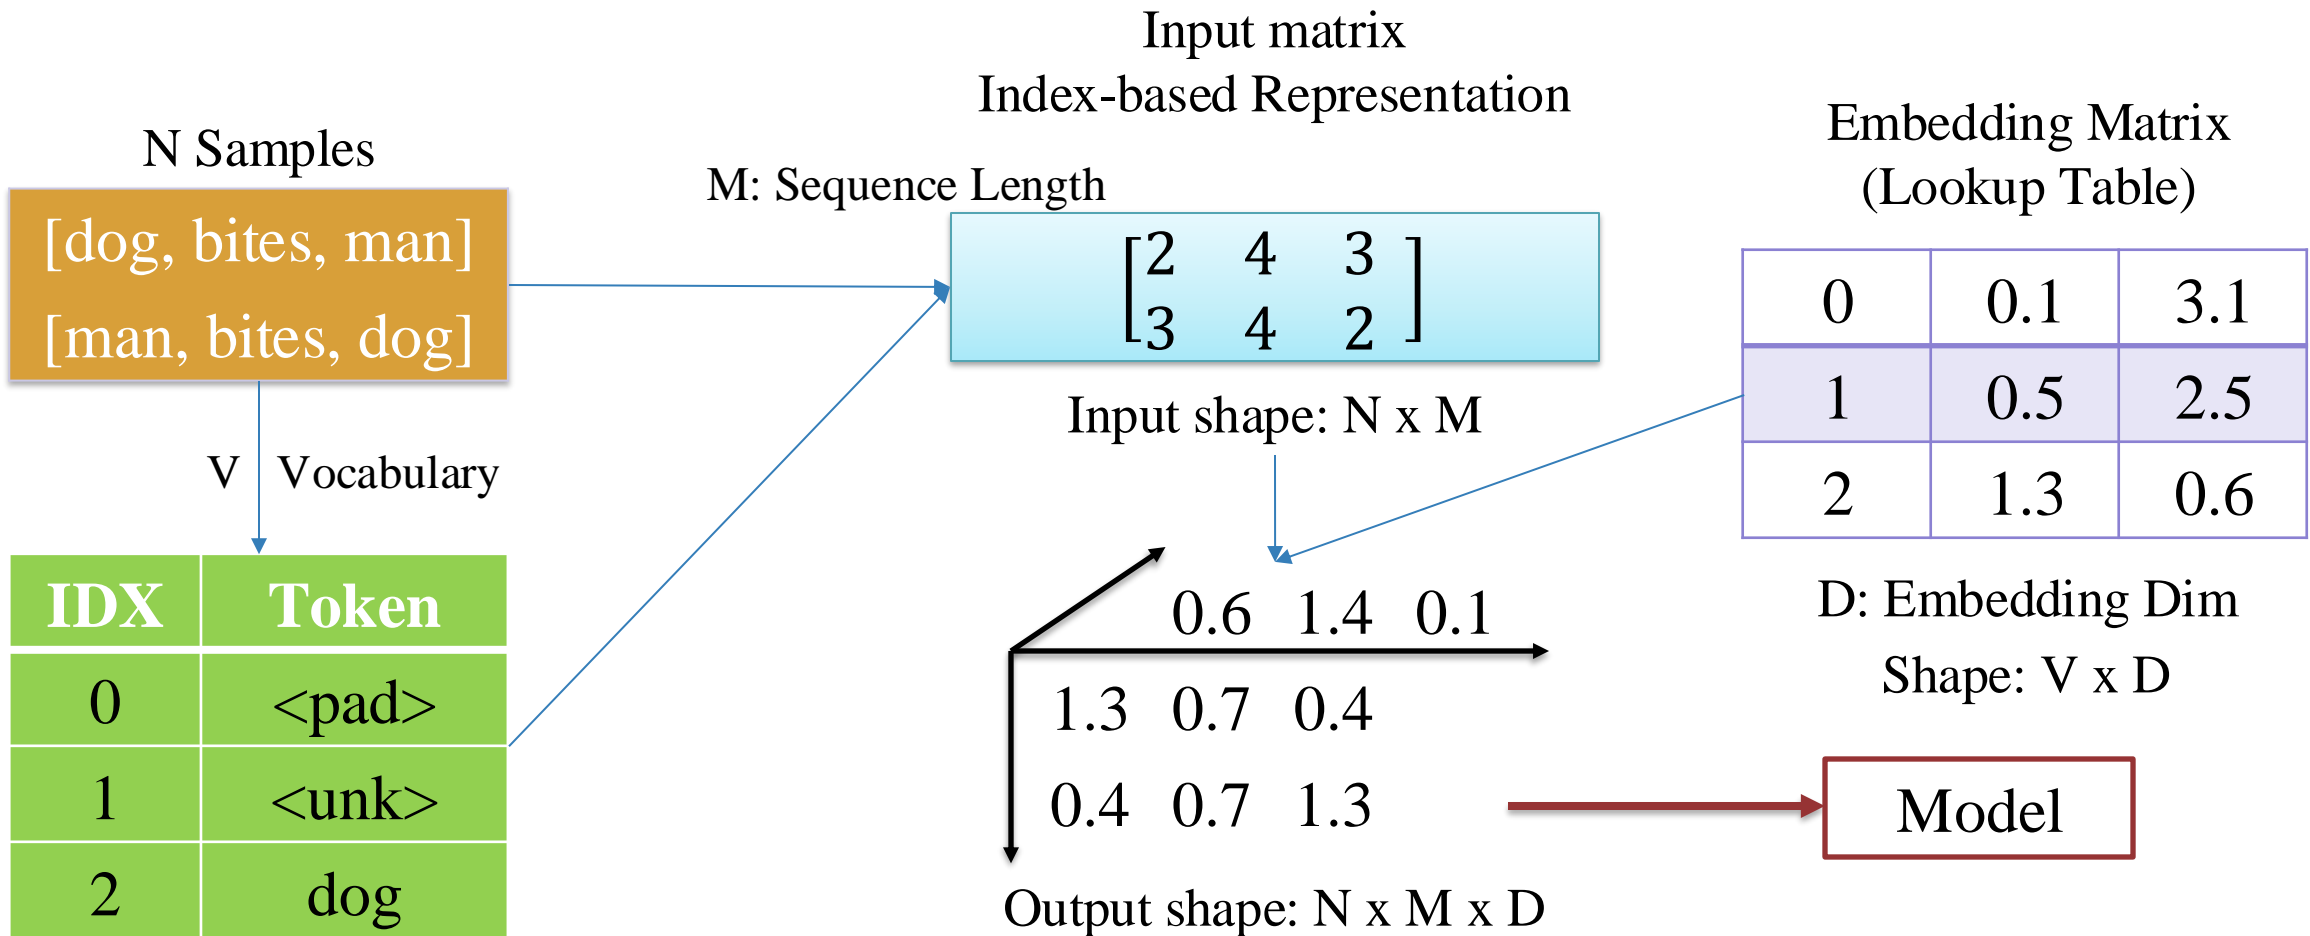
### ➤ Sentiment Analysis

Positive Example	Negative Example
Mình được 1 cô bạn giới_thiệu đến đây , tìm địa_chỉ khá dễ . Menu nước uống chất khỏi nói . Mình muốn cũng đc 8 loại nước ở đây , món nào cũng ngon và bổ_dưỡng cả .	Quán chế_biến đồ_ăn lâu , Cá_Sapa nướng ướp rất dở , sò Long ko tươi , nước_chấm ko ngon\nTôm_lại sẽ ko bao_giờ ghé nữa , ăn_dở mà uống tiền
Mỗi lần thèm trà sữa là làm 1 ly . Quán dễ kiếm , không_gian lại rộng_rải . Nhân_viên thì dễ_thương gần_gũi . Nói_chung thèm trà sữa là mình ghé Quán ở đây vì gần nhà .	Quán này thấy khá nhiều người bảo mình nên mình đã đi ăn thử , nhưng thực_sự ăn xong thấy không được như mong_đợi lắm .

# Text Classification

## ! NTC-SCV Dataset

### ➤ Data Representation



# Text Classification



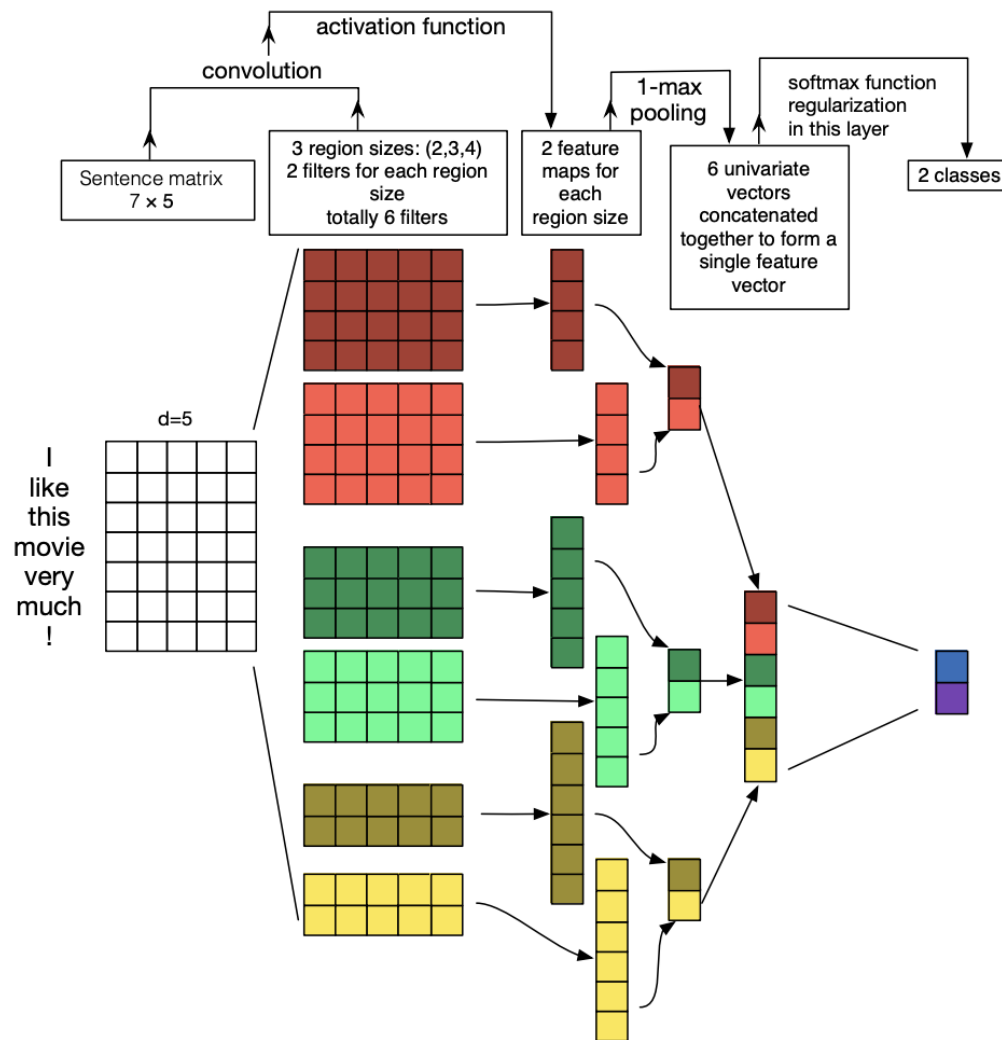
## NTC-SCV Dataset

### TextCNN Model

```
class TextCNN(nn.Module):
    def __init__(
        self,
        vocab_size, embedding_dim, kernel_sizes, num_filters, num_classes):
        super(TextCNN, self).__init__()

        self.vocab_size = vocab_size
        self.embedding_dim = embedding_dim
        self.kernel_sizes = kernel_sizes
        self.num_filters = num_filters
        self.num_classes = num_classes
        self.embedding = nn.Embedding(vocab_size, embedding_dim, padding_idx=0)
        self.conv = nn.ModuleList([
            nn.Conv1d(
                in_channels=embedding_dim,
                out_channels=num_filters,
                kernel_size=k,
                stride=1
            ) for k in kernel_sizes])
        self.fc = nn.Linear(len(kernel_sizes) * num_filters, num_classes)

    def forward(self, x):
        batch_size, sequence_length = x.shape
        x = self.embedding(x.T).transpose(1, 2)
        x = [F.relu(conv(x)) for conv in self.conv]
        x = [F.max_pool1d(c, c.size(-1)).squeeze(dim=-1) for c in x]
        x = torch.cat(x, dim=1)
        x = self.fc(x)
        return x
```







# Text Classification



## NTC-SCV Dataset

➤ Training (Test Set: 88%)



# Text Classification



## NTC-SCV Dataset

➤ Deployment (Streamlit)

### Sentiment Analysis

**Model: Text CNN. Dataset: NTC-SCV**

Sentence:

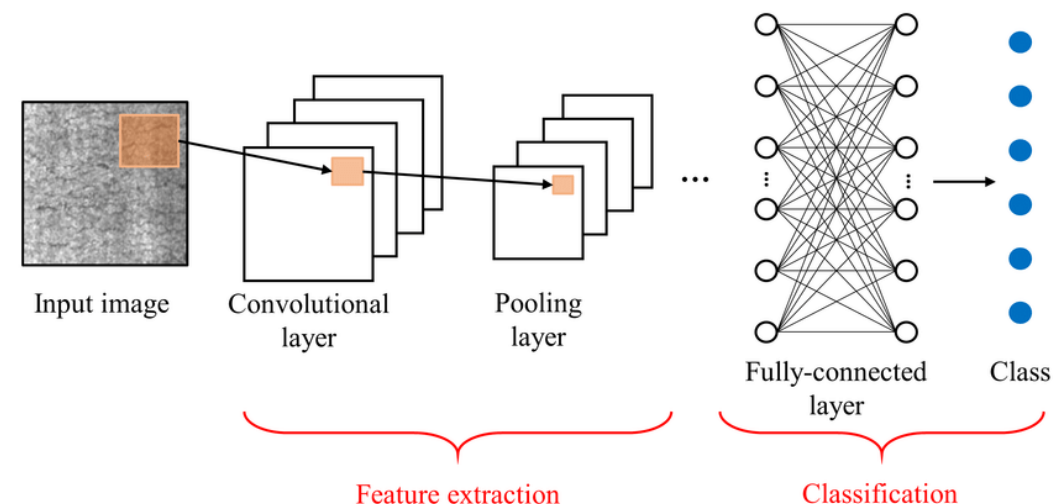

Đồ ăn ở quán này quá tệ luôn!

Sentiment: negative with 100.00 % probability.

# Objectives

## CNN

- ❖ Convolution Layer, Pooling Layer
- ❖ Multiple Input – Output Channels
- ❖ LeNet Model



## Classification

- ❖ Image Classification
- ❖ TextCNN Model
- ❖ Text Classification



AI VIET NAM

@aivietnam.edu.vn

# Thanks!

## Any questions?